Dhwani Suthar

+91-6355476267 | dhwani.suthar02@gmail.com | linkedin.com/in/dhwani-suthar

Summary

FinOps Engineer & Data Specialist with nearly 3 years of experience architecting enterprise-grade Cloud Cost Management platforms. Expert in **Cost Allocation, Showback/Chargeback models**, and Budget Forecasting across Multi-Cloud environments (GCP, AWS, Azure). Proven track record of driving $1.5M+ in annualized savings by designing and engineering custom tooling for anomaly detection and resource rightsizing. Deep technical proficiency in SQL, PySpark, and transforming complex Billing Data into actionable financial insights.

Technical Skills

FinOps & Cloud Cost: Cost Allocation Models, Chargeback/Showback, GCP Billing Exports, AWS Cost & Usage Reports (CUR), Tagging Strategies, Anomaly Detection, Budgeting & Forecasting.
Tooling & BI: Custom FinOps Dashboards, OTBI & BIP (BI Tools), SQL, Python, JIRA.
Data Engineering: Apache Spark (PySpark), Airflow, BigQuery, Databricks, ETL Pipelines.
Infrastructure: GCP, AWS, Azure, Kubernetes

Experience

Motorola Solutions July 2025 – Present
Data Engineer (FinOps & AI) *Bangalore, India*

- **Financial Modelling & Chargeback**: Architected a custom **Multi-Cloud Chargeback Platform** to allocate shared Kubernetes (GKE) and compute costs to tenant business units. Achieved 99.9% cost allocation accuracy, enabling precise showback reports for distributed teams.
- Leveraged cloud **recommendation exports** to build a centralized optimization engine, identifying $1.5M+ USD in potential savings by aggregating rightsizing metrics and idle resource data (e.g., detached volumes, underutilized VMs) into actionable dashboards.
- **Billing Data Analysis**: Engineered ETL pipelines using Python (PySpark) and SQL to ingest raw GCP Billing Exports and Prometheus metrics, transforming complex usage data into unified BI dashboards for leadership reporting.
- **Anomaly Detection**: Developed a Python-based algorithmic auditing tool to monitor daily cloud spend trends, reducing false positive budget alerts by 40% and preventing revenue leakage.
- **Governance & Tagging**: Established a Tagging Governance Framework to ensure 100% resource attribution, working with engineering teams to enforce mandatory cost-center tags via pipelines.

Mastek Jan 2024 – July 2025
Associate Consultant (Data & GenAI) *Ahmedabad, India*

- Developed a P2P (Procure-to-Pay) compliance algorithm in SQL/PLSQL that identified **£645k+** in early-receipting risks and **£214k** in unbilled spend for a major public sector client.
- Spearheaded the R&D of 'SQLator', a **Text-to-SQL** Generative AI tool using **LangChain** and **Neo4j GraphRAG**, enabling non-technical stakeholders to query complex Oracle Finance schemas via natural language.
- Designed and deployed **Oracle APEX** dashboards to visualize OEE (Overall Equipment Effectiveness) metrics across three manufacturing plants, integrating IoT data streams to minimize production downtime.
- Automated the ETL process for "Certificate of Analysis" generation, reducing manual reporting effort by **90%** through dynamic SQL templating and stored procedures.

Jio Platforms Ltd. May 2023 – July 2023
Data Engineering Intern (Observability) *Mumbai, India*

- Processed high-volume system logs using **Splunk** (SPL), building real-time dashboards to monitor distributed system health and API error rates.
- Developed Python scripts to parse unstructured log data into structured formats for downstream anomaly analysis.

Key Projects

Autonomous FinOps Agent Swarm | *LangGraph, Vertex AI, BigQuery*

- **Orchestrated a Multi-Agent System (MAS)** using LangGraph and Gemini, employing a Supervisor-Worker architecture to route tasks between specialized agents (Analyst vs. Optimizer) for autonomous cloud cost analysis.
- **Engineered self-correcting SQL generation** capable of dynamic schema discovery and validation, enabling the agent to write complex BigQuery SQL from natural language without hallucinations.
- **Developed a hybrid reasoning engine** that chains SQL retrieval with Python-based logic tools to simulate **Committed Use Discount (CUD)** scenarios, providing real-time savings estimates and interactive Streamlit visualizations.

Scalable Azure Data Lake Pipeline | *PySpark, Databricks, Delta Lake, Unity Catalog*

- Built a production-grade ETL pipeline using the **Medallion Architecture** (Bronze/Silver/Gold) on Azure Databricks.
- Implemented **SCD Type 1 & 2** logic for dimension tables and incremental watermarking for fact tables, optimizing compute costs.
- Enforced data governance and fine-grained access control (RBAC) using **Unity Catalog**.

Predictive Maintenance with AutoEncoders | *Vertex AI, TensorFlow, BigQuery ML*

- Developed an unsupervised **Deep Learning AutoEncoder** model on Vertex AI to detect anomalies in turbofan engine sensor data.
- Deployed the model endpoint on GCP, utilizing **reconstruction error thresholds** to flag potential hardware failures before they occurred.

GraphRAG Knowledge Engine | *Neo4j, LangChain, LLMs*

- Built a Retrieval-Augmented Generation (RAG) system leveraging Knowledge Graphs to map relationships between 50+ database tables, significantly improving LLM hallucination rates on schema queries.

Education

Charotar University of Science & Technology May 2024
Bachelor of Technology in Computer Engineering *CGPA:* ***9.79 / 10.0***

Awards & Certifications

Certifications:

- FinOps Foundation: Introduction to FinOps
- Apache Airflow 3 Fundamentals (Astronomer)
- Oracle Cloud GenAI Professional

Awards: "People First Growth Award" (Motorola Solutions), "The Dream Team" & "A True Ideator" (Mastek).

Hackathon: State-Level Winner, Azadi Ka Amrit Mohotsav Hackathon.